

# Stellar-Superccluster

an automation tool for stellar-core

# First things first!

- `git clone https://github.com/stellar/stellar-supercluster`
- Install an F# mode in your editor if you like
- Follow along!

# Talk Overview

- What SSC is intended for vs. not-intended for
- What it depends on to build and run
- Crash course in Kubernetes
- How SSC uses Kubernetes
- How the SSC codebase is organized
- How to run it and debug it
- How to extend it

# Part 1:

## What is it intended for (and not-intended for)

# Primary purpose: replace SCC (with SSC, they're different!)

- Long-standing automation tool: `stellar_core_commander`
  - Written in Ruby, runs `stellar-core` locally or via Docker passing lots of environment variables
  - `confd` inside Docker image expands config templates using environment variables
  - Shell script entry-point in Docker image to wrap core tasks

# What was wrong with SCC?

- Very hard to debug, extend, or even keep working
  - Too many steps to edit, debug, version, upgrade, make mistakes in: Ruby DSL → Ruby library → Env vars → Dockerfile → Shell script → Confd → stellar-core
  - No typechecking or even typo-checking. Easy to modify an hours-long test and have it fail right at the end because of a string-pasting config error
- Limited to small set of Docker hosts, hard to scale or redeploy
- No plan for how to extend to work with Kubernetes

# How does SSC improve on things?

- Types: more coding errors caught before runtime, can use IDE / autocomplete
  - Stellar SDK is strongly typed
  - Kubernetes API is strongly typed
  - Internal abstractions in SSC (configs, HTTP JSON responses) all typed
- Kubernetes has more visibility and debugging tools, robust controls
  - All configs, logs, status of all components viewable by `kubectl` commands
  - Kubernetes reports metrics to Prometheus, easy to observe
  - CPU / RAM quotas enforced, much harder to crash the cluster
- Can scale cluster arbitrarily by adding more workers, scheduling is dynamic
  - Single developer gets remote access to \$LOTS\_OF\_CORES

# Ok but what was SCC intended for?

- Acceptance tests
  - Run by Jenkins "as often as possible", typically ~daily
  - Bigger & slower than merge-gating CI unit tests
- Performance / load tests
  - Run by hand on developer workstation
  - Evaluate overall scalability of core



# Acceptance Tests

- Usually several **stellar-core** instances, communicating
- Either simulated private networks or the real live networks
- Version mixtures, protocol and database upgrades
- Generation of history archives, replay of history archives
- Most thorough test: replay all of pubnet's history

# What is not intended

- Not for managing live operational deployments
  - If you hit ctrl-c while it's running, it has no good way of recovering aside from "delete all the Kubernetes objects in the namespace"
- Not for small iterative unit tests while developing
  - You have to have a Docker image to run, which only happens after CI passes

# Part 2:

## What it depends on to build and run

# Software Dependencies

- Stellar-Supercluster (SSC) is written in F#
  - This is a .NET language, so requires .NET Core which is available on Linux, MacOS and Windows
  - Build in root dir with `dotnet build`
  - The package is split in two pieces, one C#, one F#, so you can add C# parts if you really hate F#, but I like it

# Software Dependencies

- Uses several libraries, all auto-installed when you build
  - The C# Kubernetes client
  - The C# Stellar SDK
  - The FSharp.Data JSON and CSV data-access library
  - The Nett TOML library
  - The Serilog logging library
  - The CommandLineParser library

# Runtime Dependencies

- SSC runs against a Kubernetes cluster only
  - You need a `kubeconfig` file that lets you talk to one
- It runs a Docker image on the cluster
  - You need one that contains `stellar-core`
  - The default is `stellar/stellar-core` which is ok if you are happy tracking master
- Of course, F#/C# means you need the `dotnet` runtime installed on the client you run SSC from

# Aside: why F#?

- It's strongly-typed but type-inferred and relatively terse.
  - "Missions" (a.k.a. "recipes" in SCC a.k.a. "tests") can be written directly in it without much boilerplate; it's "terse as a DSL".
- Immutable-by-default, polymorphic, functional style: relatively high expressivity and safety.
- .NET ecosystem: lots of well-typed libraries, cross-platform, decent JIT'ed performance when generating load / ingesting / scanning large output.

# Example mission

```
let simplePayment (context: MissionContext) =  
  let coreSet = MakeLiveCoreSet "core" CoreSetOptions.Default  
  context.Execute [coreSet] None (fun (formation: StellarFormation) ->  
    formation.WaitUntilSynced [coreSet]  
    formation.UpgradeProtocolToLatest [coreSet]  
  
    formation.CreateAccount coreSet UserAlice  
    formation.CreateAccount coreSet UserBob  
    formation.Pay coreSet UserAlice UserBob  
  )
```



# Part 3:

# Crash course in

# Kubernetes

# Kubernetes Basics

- Resources: things client gets to talk about: get, put, list, patch
  - Volumes, Containers, Pods (groups-of-containers), Jobs, StatefulSets (groups-of-pods), Services, Ingresses ...
  - Each resource has a State and a Spec
  - Client can observe State, but can only modify Spec
- Controllers: server-side processes that modify resource States
  - DNS Controller, Ingress Controller, Job Controller, PersistentVolume Controller ...
  - Observe changes to resource Specs made by clients
  - Try to make State match Spec, if possible
    - Possibly instantiating more Resources, contacting other Controllers

# Kubernetes Basics

- All resources have uniform YAML representation
  - `kubectl list` / `kubectl get` / `kubectl describe`
- Can dump to YAML file, edit, patch back into server
  - We do not do this except while debugging!
- Also exposed as uniform REST interface, mapped to strongly typed client SDKs for navigating & updating
  - We do this!

# Kubernetes Basics

- SSC interaction pattern:
  - Connect to Kubernetes API server
  - Build client-side big pile of Resource Spec objects
  - Send REST requests to **create** all such Resources
  - Make long-poll **watch** requests until States match
  - Send HTTP commands to **stellar-core** to observe & inject traffic
    - Ingress Resource routes URL prefix to each container
  - Possibly **replace** Spec to reflect next test phase, re-**watch**

# Kubernetes Unique Terms

- Pod: bundle of two sets-of-resources
  - Set of Containers running together on a node
  - Set of shared Volumes visible to the Containers
- Example (single Pod):
  - `stellar-core` container + `nginx` history-serving container
  - ConfigMap Volume + EmptyDir or EBS data Volume

# Kubernetes Unique Terms

- StatefulSet: group of similar pods
  - PodTemplate that gets replicated for i in 0..N-1
  - Internal DNS names assigned
    - `$podname-$i.$service.$namespace.svc.cluster.local`
  - EBS volumes acquired from provisioner, attached
- Example (single StatefulSet):
  - Quorum of N `stellar-core` + `nginx` pods
  - Each configured to see each other in own `/cfg/$podname-$i.cfg`:
    - `PREFERRED_PEERS=["$podname-0.$service...", "$podname-1.service...", ...]`
    - `[HISTORY.$podname-0]`  
`get="curl -sf http://$podname-0.$service.../{0} -o {1}"`  
  
`[HISTORY.$podname-1]`  
`get="curl -sf http://$podname-1.$service.../{0} -o {1}"`

# Kubernetes Namespaces

- All SSC Resources created in a Namespace
  - Umbrella Resource for holding other Resources
  - Access control, name isolation, resource quotas
  - SSC Resource names further randomly prefixed
- Namespace typically cleared pre-run by Jenkins
  - Get your own (ask admins) if you want isolated workspace for personal use!

# Kubernetes Quotas

- RAM and CPU subject to Quotas
  - Admins set the quotas, we obey.
  - Namespace has overall Quotas; Containers have sub-quotas within.
- Quotas contain both a Request (guaranteed) and Limit (opportunistic) value. Must not allocate more of either than Quota allows, or will not be scheduled to run.
- Calculation of the "right" amount is complicated!
  - As much as possible to make room; not so much that Quota exceeded.
  - Depends on how many containers you run, how many SSC instances, etc.
  - SSC does these calculations, has many knobs to adjust.



# Part 4:

# How SSC uses

# Kubernetes

# Resources Used

- Containers
  - To run `stellar-core`, `nginx`, `sqlite3`, maybe `postgresql`
- Volumes
  - ConfigMaps to store TOML `.cfg` files
  - EBS or EmptyDirs to store data (buckets, DB, history)
  - EBS Volumes use PersistentVolumeClaims dynamically provisioned by admin-configured "storage class"

# Resources Used

- Pods and PodTemplates
  - Combine Volumes and Containers
  - Pod can have 1 or more Initialization commands, eg. `stellar-core new-db`
  - PodTemplate instantiated into multiple Pods by StatefulSet

# Resources Used

- StatefulSets and Jobs
  - StatefulSet for group of stellar-core pods with ongoing overlay connections among themselves, eg. a quorum
  - Job for isolated run-then-exit commands such as parallel catchup tests that run `stellar-core catchup`

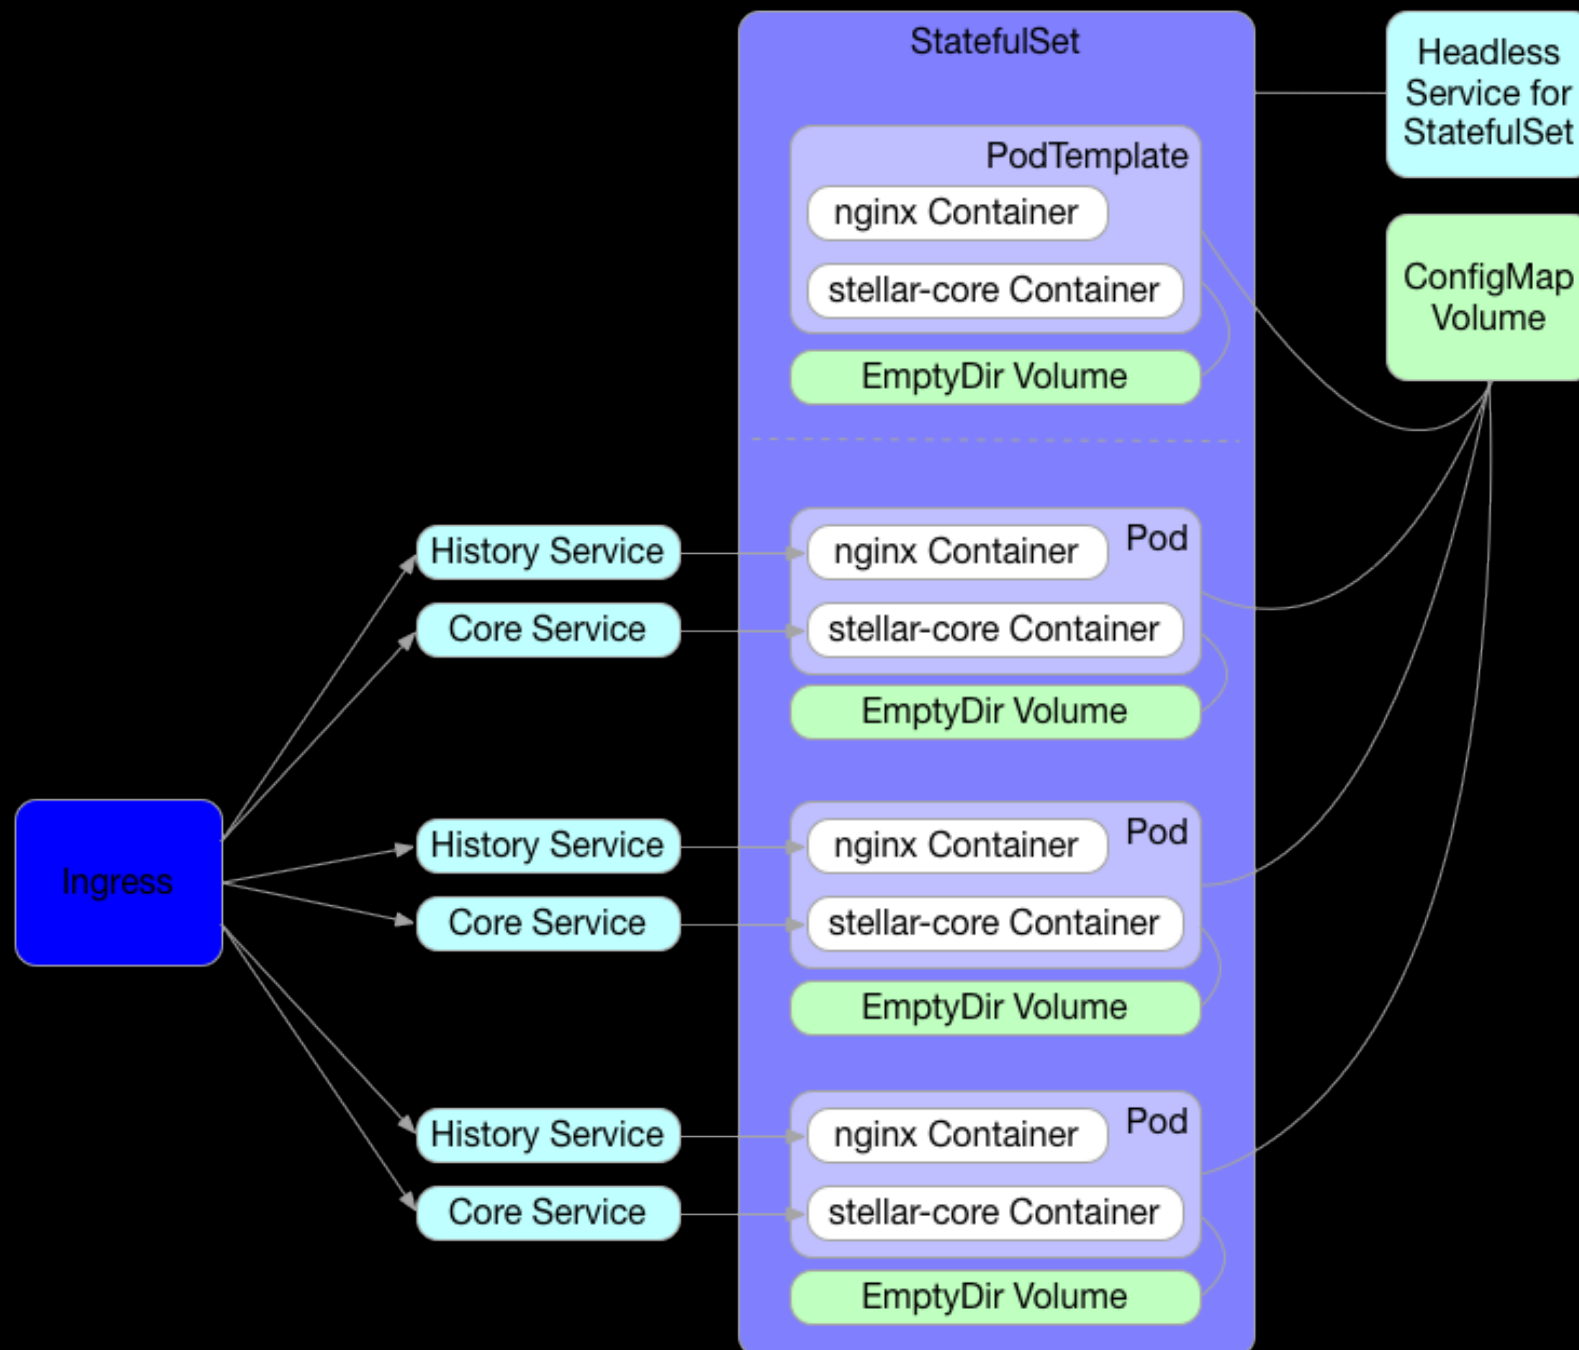
# Resources Used

- Services
  - Specifies network port Ingress might route traffic to
  - Can be used for load balancing, but we do not
  - We make two Services per-Pod in each StatefulSet:
    - History Service to talk to the Pod's **nginx**
    - Core Service to talk to the Pod's **stellar-core**

# Resources Used

- Ingress
  - Listens on port at cluster edge
  - Runs **traefik** "cloud-native edge router" (HTTP proxy)
  - Routes HTTP requests to (per-Pod) Services inside:
    - History: **`http://$ingress/$peer-$i/history/...`**
    - Core: **`http://$ingress/$peer-$i/core/...`**

# Diagram Time!



# Part 5:

## How the SSC codebase is organized



# Root directory

- **Dockerfile.dotnet**
  - Makes a build environment containing **dotnet**, stores result as Docker image
- **Dockerfile**
  - Runs **dotnet build StellarSupercluster.sln**, stores result as Docker image
- **StellarSupercluster.sln**
  - "Solution" file for **dotnet**, references sub-projects in **src/**
- **src/\*\***
  - The code
- **Jenkinsfile**
  - Groovy code, builds Docker images above, runs acceptance tests

# src directory

- **App**
  - Command-line parsing and sub-command definitions
- **CSLibrary**
  - C# sub-project, currently an empty stub
- **FSLibrary**
  - F# sub-project, currently almost all of the code

# FSLibrary directory

- **Mission\*.fs**
  - Specific acceptance-test scenarios ("missions"), one per file.
- **Stellar\*.fs**
  - Logic to write stellar-core configs, build Kubernetes Specs, watch Kubernetes States, build Stellar SDK transactions, speak HTTP to Ingress
- **{csv,json}-type-samples/\*\***
  - Sample JSON and CSV data from which static types are derived

# Primary Types

- `StellarCoreCfg.fs` (type `StellarCoreCfg`)
  - Immutable config for single stellar-core node (with TOML writing)
- `StellarCoreSet.fs` (type `CoreSet`)
  - Immutable set of related stellar-core config values (eg. a quorum)
- `StellarNetworkCfg.fs` (type `NetworkCfg`)
  - Immutable set of CoreSets and collective network config (Quota, Ingress, Namespace, nonce, network passphrase)
- `StellarFormation.fs` (type `StellarFormation`)
  - Mutable NetworkCfg and Kubernetes objects

# Secondary Types

- **StellarCorePeer.fs (type Peer)**
  - Handle to a specific member of a CoreSet within a NetworkCfg
- **StellarDestination.fs (type Destination)**
  - Client-side directory in which logs & data dumps written
- **StellarPerformanceReporter.fs (type PerformanceReporter)**
  - Accumulates sequence of metrics from Peer, writes CSV
- **StellarNamespaceContent.fs (type NamespaceContent)**
  - Tracks live Kubernetes Resources, deletes them on shutdown

# Type Extensions

- `StellarCoreHTTP.fs` (extends type `Peer`)
  - Adds methods to speak HTTP to stellar-core through Ingress
- `StellarTransaction.fs` (extends type `Peer`)
  - Adds methods to compose XDR / Stellar SDK-level objects to send over HTTP
- `StellarDataDump.fs` (extends type `StellarFormation`)
  - Adds methods to extract logs, SQL DB, etc. to Destination
- `StellarJobExec.fs` (extends type `StellarFormation`)
  - Adds methods to run one-off Jobs (including in parallel)
- `StellarKubeSpecs.fs` (extends type `NetworkCfg`)
  - Adds methods to translate `NetworkCfg` content to Kubernetes Specs

# Part 6:

## How to run it and debug it

# Running

- This part is a little clunky. Fundamentally just:
  - `dotnet run ...`
- But in practice, lots of args:
  - `dotnet run --project src/App/App.fsproj  
--configuration Release  
--  
mission SomeMissionName  
--kubeconfig path/to/kubeconfig  
--destination path/to/destination  
--image stellar/stellar-core:latest  
--old-image stellar/stellar-core:$vers`



# Debugging

- Mainly you want:
  - `kubectl logs <pod-name>`
  - `kubectl (get|list|describe) <resource-type> <name>`
  - `kubectl (exec|cp|port-forward) ...`
- Also:
  - Watch Prometheus metrics
  - Browse to the Kubernetes dashboard (if you're allowed)
  - Run with `--keep-data` and inspect Resources, then manually delete with `kubectl delete --all pod,svc,sts,ing,pvc,cm,job`

# Part 7:

## How to extend it

# Extensions: missions

- Add new missions:
  - Duplicate and modify `src/FSLibrary/Mission*.fs` file
- Add any parameters needed
  - Add to `src/App/Program.fs` arg-parsing code
  - Thread through to `src/FSLibrary/MissionContext.fs`
  - From there to wherever else they're needed
- Wire in to `Jenkinsfile`, adjusting `--num-concurrent-missions`

# Extensions: Resources

- Customize existing Kubernetes Resources:
  - Edit Spec-forming code in `StellarKubeSpecs.fs`
- Add new Kubernetes Resources:
  - Add methods to make Specs in `StellarKubeSpecs.fs`
  - Create Resources in `Kubernetes.MakeFormation` extension method in `StellarSupercluster.fs`

# Extensions: protocol

- Add HTTP endpoint calls to stellar-core:
  - Add methods to `StellarCoreHTTP.fs`
- Add XDR / SDK-typed interactions with stellar-core:
  - Add methods to `StellarTransaction.fs`

**Fini**